

Machine Learning Worksheet 2

Tomas Ladek
3602673
tom.ladek@tum.de

Problem 1

After parsing the data that was given in a csv file and making up an efficient data structure (matrix), the Gini index of the root node ($C = \{0, 1, 2\}$) was calculated:

$$i_G(t) = 1 - \left(\frac{5}{15}\right)^2 - \left(\frac{6}{15}\right)^2 - \left(\frac{4}{15}\right)^2$$

Then in 0.1 steps from -0.6 to +10.0 (limits determined by data inspection), the Gini index of all possible left/right splits of the root node was calculated, for all features $x_{i,1} \dots x_{i,3}$. The maximum was a difference in Gini indices of ≈ 0.3615 for the first feature ($x_{i,1}$) for a split at the value 4.5¹. The formula used for calculating the difference was

$$\Delta i_G(t)(x_{i,1} \leq s, t) = i_G(t) - \frac{\text{\#classes in left tree}}{\text{\#classes in current data set}} i_G(t_L) - \frac{\text{\#classes in right tree}}{\text{\#classes in current data set}} i_G(t_R)$$

with i_G being the Gini indices of the current node, the left tree and the right tree respectively.

Splitting the root node at $x_{i,1} = 4.5$ yielded a left tree (values less than or equal to the split value) consisting of a pure node (class distribution of 100% for class '1', Gini index 0) and a right tree combining the remaining classes '0' and '2', Gini index 0.4938. No further splits in the left tree were needed.

Once again performing a maximalization on the Gini index differences for every possible split of every feature in 0.1 steps yielded feature $x_{i,1}$ at 7.4 as the best split. The result was a left sub-tree with class distribution '2': $\frac{2}{3}$; '0': $\frac{1}{3}$ and a right sub-tree with a pure distribution of class '0'. The corresponding Gini indices were 0.4444 and 0 respectively.

As the maximum requested depth was reached, no further splitting was considered. The (raw) python code can be provided on request.

Problem 2

By using the tree that was constructed and explained above, the vector \mathbf{x}_a is classified as '1' (following the left sub-tree, since $\mathbf{x}_{a,1} \leq 4.5$). Being a pure leaf node, $p(c = 1 | \mathbf{x}_a, T) = 1$.

Vector \mathbf{x}_b follows the right sub-tree ($\mathbf{x}_{b,1} > 4.5$) and ends up in its left leaf ($(\mathbf{x}_{b,1} \leq 7.4)$). The corresponding classification is '2', because '2' is the majority class in that node. The probability of a correct classification is $p(c = 2 | \mathbf{x}_b, T) = \frac{2}{3}$, due to the class distribution of that node.

¹Due to the chosen calculation procedure, the split values take the maximum possible value between two different splits, instead of the average.

Problem 3

02_homework_knn

November 14, 2016

```
In [1]: import random
import numpy as np
import operator
from sklearn import datasets
import matplotlib.pyplot as plt
%matplotlib inline
```

0.1 Load dataset

The iris data set (https://en.wikipedia.org/wiki/Iris_flower_data_set) it loaded by the function loadDataset.

Arguments:

- *split*: int: Split rate between test and training set e.g. 0.67 corresponds to 1/3 test and 2/3 validation

Returns:

- X: list(array of length 4); Trainig data
- Z: list(int); Training labels
- XT: list(array of length 4); Test data
- ZT: list(int); Test labels

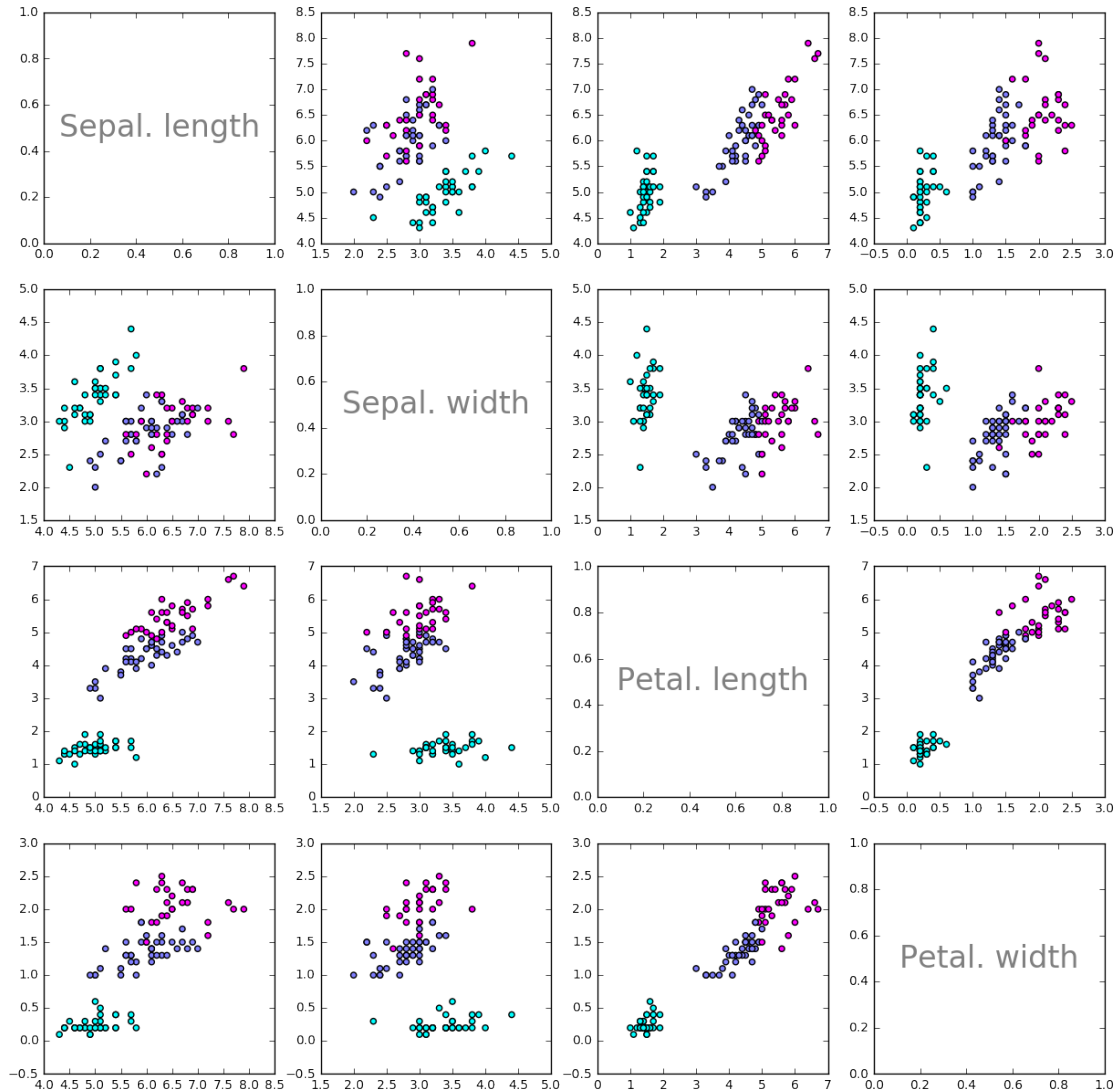
```
In [2]: def loadDataset(split, X=[], XT=[], Z = [], ZT = []):
dataset = datasets.load_iris()
c = list(zip(dataset['data'], dataset['target']))
random.seed(224)
random.shuffle(c)
x, t = zip(*c)
sp = int(split*len(c))
X = x[:sp]
XT = x[sp:]
Z = t[:sp]
ZT = t[sp:]
return X, XT, Z, ZT
```

```
In [3]: # prepare data
split = 0.67
X, XT, Z, ZT = loadDataset(split)
```

0.2 Plot dataset

Since X is dimensionality 4, 16 scatterplots (4x4) are plotted showing the dependencies of each two features.

```
In [4]: Xa = np.asarray(X)
        f, axes = plt.subplots(4, 4, figsize=(15, 15))
        for i in range(4):
            for j in range(4):
                if j == 0 and i == 0:
                    axes[i,j].text(0.5, 0.5, 'Sepal. length', ha='center', va='center')
                elif j == 1 and i == 1:
                    axes[i,j].text(0.5, 0.5, 'Sepal. width', ha='center', va='center')
                elif j == 2 and i == 2:
                    axes[i,j].text(0.5, 0.5, 'Petal. length', ha='center', va='center')
                elif j == 3 and i == 3:
                    axes[i,j].text(0.5, 0.5, 'Petal. width', ha='center', va='center')
                else:
                    axes[i,j].scatter(Xa[:,j],Xa[:,i], c = Z, cmap=plt.cm.cool)
```



0.3 Exercise 1: Euclidean distance

Compute euclidean distance between two data points.

arguments: * $x1$: array of length 4; data point * $x2$: array of length 4; data point

returns: * $distance$: float; euclidean distance between $x1$ and $x2$

1 Implementation exercise: k-NN

```
In [5]: def euclideanDistance(x1, x2):
        return math.sqrt( pow((x1[0] - x2[0]), 2) + pow((x1[1] - x2[1]), 2) + pow((x1[2] - x2[2]), 2) + pow((x1[3] - x2[3]), 2) )
```

1.1 Exercise 2: get k nearest neighbors

For one data point x_t compute all k nearest neighbors.

arguments: * X : list(array of length 4); Trainig data * Z : list(int); Training labels * x_t : array of length 4; Test data point

returns: * neighbors: list of length k of tuples (X_{neighbor} , Z_{neighbor} , distance between neighbor and x_t); **this is the list of k nearest neighbors to x_t**

```
In [6]: def getNeighbors(X, Z, xt, k):
        Xnp = np.array(X)
        datapoints = Xnp.shape[0]
        distance = np.empty(datapoints)
        for i in range(datapoints):
            distance[i] = euclideanDistance(Xnp[i], xt)
        nearest = distance.argsort()[:k]

        nearest_list.append((X[i], Z[i], distance[i]))
        return nearest_list
```

1.2 Exercise 3: get neighbor response

For the previously computed k nearest neighbors compute the actual response. I.e. give back the class of the majority of nearest neighbors. What do you do with a tie?

arguments: * neighbors: list((array, int, float) * c : int; number of classes in the dataset, for the iris dataset $c=3$

returns * y : int; majority target

```
In [7]: def getResponse(neighbors, c=3):
        classes = []
        labelCounts = list([0, 0, 0])
        for neighbor in neighbors:
            classes.append(neighbor[1])
        classes = list(set(classes))
        for neighbor in neighbors:
            if neighbor[1] == classes[0]:
                labelCounts[0] = labelCounts[0] + 1
            elif neighbor[1] == classes[1]:
                labelCounts[1] = labelCounts[1] + 1
            elif neighbor[1] == classes[2]:
                labelCounts[2] = labelCounts[2] + 1
        classIndex = labelCounts.index(max(labelCounts))
        return classes[classIndex]
```

1.3 Exercise 4: Compute accuracy

arguments: * YT : list(int); predicted targets * ZT : list(int); actual targets

returns: * accuracy: float; percentage of correctly classified test data points

```
In [8]: def getAccuracy(YT, ZT):
        return 0
```

```
In [9]: def predict(X, Z, XT, k):
        Y=[]
        for xt in XT:
            neighbors = getNeighbors(X, Z, xt, k)
            Y.append(getResponse(neighbors))
        return Y
```

1.4 Testing

Should output an accuracy of 0.95999999999999996.

```
In [10]: # prepare data
        split = 0.67
        X, XT, Z, ZT = loadDataset(split)
        print 'Train set: ' + repr(len(X))
        print 'Test set: ' + repr(len(XT))
        # generate predictions
        k = 3
        YT = predict(X, Z, XT, k)
        accuracy = getAccuracy(YT, ZT)
        print('Accuracy: ' + repr(accuracy))
```

```
File "<ipython-input-10-4285cf12047f>", line 4
print 'Train set: ' + repr(len(X))
      ^
```

SyntaxError: invalid syntax

```
In [ ]:
```

Problem 4

Euclidean distance: $d = \sqrt{\sum_i (u_i - v_i)^2}$

Distances for \mathbf{x}_a :

$$d_A = 2.76 \quad d_B = 3.78 \quad d_C = 2.18 \quad d_D = 5.97 \quad d_E = 3.22 \quad d_F = 3.97 \quad d_G = 2.93 \quad d_H = 2.83$$

$$d_I = 0.67 \quad d_J = 3.6 \quad d_K = 2.84 \quad d_L = 5.3 \quad d_M = 3.16 \quad d_N = 4.26 \quad d_O = 2.47$$

\Rightarrow 3-nearest neighbors are I, C, O corresponding to classes ('0', '2' and '1' respectively). Therefore:

$$p(z = 0|x, 3) = \frac{1}{3} \cdot 1$$

$$p(z = 1|x, 3) = \frac{1}{3} \cdot 1$$

$$p(z = 2|x, 3) = \frac{1}{3} \cdot 1$$

Distances for \mathbf{x}_b :

$$d_A = 3.26 \quad d_B = 2.73 \quad d_C = 2.12 \quad d_D = 3.84 \quad d_E = 1.17 \quad d_F = 3.93 \quad d_G = 4.3 \quad d_H = 4.86$$

$$d_I = 1.75 \quad d_J = 5.7 \quad d_K = 3.15 \quad d_L = 3.77 \quad d_M = 5.32 \quad d_N = 6.45 \quad d_O = 4.56$$

\Rightarrow 3-nearest neighbors are E, I, C corresponding to classes ('2', '0' and '2' respectively). Therefore:

$$p(z = 0|x, 3) = \frac{1}{3} \cdot 1 = \frac{1}{3}$$

$$p(z = 2|x, 3) = \frac{1}{3} \cdot 2 = \frac{2}{3}$$

Problem 5

$$y = \frac{1}{\sum_{i \in N_k(x)} \frac{1}{d(x, x_i)}} \sum_{i \in N_k(x)} \frac{1}{d(x, x_i)} z_i$$

For \mathbf{x}_a : $y = 0.56$

For \mathbf{x}_b : $y = 1.398$